

# Advanced Tabular Data

ACTL3143 & ACTL5111 Deep Learning for Actuaries  
Patrick Laub



# Lecture Outline

- **Entity Embedding**
- Categorical Variables & Entity Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale By Exposure
- Dropout



# Continuing on the French motor dataset example

Download the dataset if we don't have it already.

```
1 from pathlib import Path
2 from sklearn.datasets import fetch_openml
3
4 if not Path("french-motor.csv").exists():
5     freq = fetch_openml(data_id=41214, as_frame=True).frame
6     freq.to_csv("french-motor.csv", index=False)
7 else:
8     freq = pd.read_csv("french-motor.csv")
9
10 freq
```



Source: Nell et al. (2020), [Case Study: French Motor Third-Party Liability Claims](#), SSRN.

# Continuing on the French motor dataset example

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAg
0	1.0	1	0.10000	D	5	0
1	3.0	1	0.77000	D	5	0
...	...	...	...	...	...	...
678011	6114329.0	0	0.00274	B	4	0
678012	6114330.0	0	0.00274	B	7	6

678013 rows  $\times$  12 columns

# Data dictionary

- **IDpol**: policy number (unique identifier)
- **ClaimNb**: number of claims on the given policy
- **Exposure**: total exposure in yearly units
- **Area**: area code (categorical, ordinal)
- **VehPower**: power of the car (categorical, ordinal)
- **VehAge**: age of the car in years
- **DrivAge**: age of the (most common) driver in years
- **BonusMalus**: bonus-malus level between 50 and 230 (with reference level 100)
- **VehBrand**: car brand (categorical, nominal)
- **VehGas**: diesel or regular fuel car (binary)
- **Density**: density of inhabitants per km<sup>2</sup> in the city of the living place of the driver
- **Region**: regions in France (prior to 2016)



# The model

Have  $\{(\mathbf{x}_i, y_i)\}_{i=1, \dots, n}$  for  $\mathbf{x}_i \in \mathbb{R}^{47}$  and  $y_i \in \mathbb{N}_0$ .

Assume the distribution

$$Y_i \sim \text{Poisson}(\lambda(\mathbf{x}_i))$$

We have  $\mathbb{E}Y_i = \lambda(\mathbf{x}_i)$ . The NN takes  $\mathbf{x}_i$  & predicts  $\mathbb{E}Y_i$ .

## Note

For insurance, *this is a bit weird*. The exposures are different for each policy.

$\lambda(\mathbf{x}_i)$  is the expected number of claims for the duration of policy  $i$ 's contract.

Normally,  $\text{Exposure}_i \notin \mathbf{x}_i$ , and  $\lambda(\mathbf{x}_i)$  is the expected rate *per year*, then

$$Y_i \sim \text{Poisson}(\text{Exposure}_i \times \lambda(\mathbf{x}_i)).$$

# Where are things defined?

In Keras, string options are used for convenience to reference specific functions or settings.

```
1 model = Sequential([
2     Dense(30, activation="relu"),
3     Dense(1, activation="exponential")
4 ])
```

is the same as

```
1 from keras.activations import relu, exponential
2
3 model = Sequential([
4     Dense(30, activation=relu),
5     Dense(1, activation=exponential)
6 ])
```

```
1 x = [-1.0, 0.0, 1.0]
2 print(relu(x))
3 print(exponential(x))
```

```
tf.Tensor([0. 0. 1.], shape=(3,), dtype=float32)
tf.Tensor([0.37 1.    2.72], shape=(3,), dtype=float32)
```



# String arguments to `.compile`

When we run

```
1 model.compile(optimizer="adam", loss="poisson")
```

it is equivalent to

```
1 from keras.losses import poisson
2 from keras.optimizers import Adam
3
4 model.compile(optimizer=Adam(), loss=poisson)
```

Why do this manually? To adjust the object:

```
1 optimizer = Adam(learning_rate=0.01)
2 model.compile(optimizer=optimizer, loss="poisson")
```

or to get help.





# Keras' “poisson” loss

```
1 help(keras.losses.poisson)
```

Help on function poisson in module keras.src.losses.losses:

```
poisson(y_true, y_pred)
```

Computes the Poisson loss between y\_true and y\_pred.

Formula:

```
```python
loss = y_pred - y_true * log(y_pred)
```
```

Args:

y\_true: Ground truth values. shape = `[batch\_size, d0, .. dN]`.

y\_pred: The predicted values. shape = `[batch\_size, d0, .. dN]`.

Returns:

Poisson loss values with shape = `[batch\_size, d0, .. dN-1]`.

Example:



# Subsample and split

```
1 freq = freq.drop("IDpol", axis=1).head(25_000)
2
3 X_train, X_test, y_train, y_test = train_test_split(
4     freq.drop("ClaimNb", axis=1), freq["ClaimNb"], random_state=2023)
5
6 # Reset each index to start at 0 again.
7 X_train = X_train.reset_index(drop=True)
8 X_test = X_test.reset_index(drop=True)
```



# What values do we see in the data?

```
1 X_train["Area"].value_counts()
2 X_train["VehBrand"].value_counts()
3 X_train["VehGas"].value_counts()
4 X_train["Region"].value_counts()
```

Area

C 5507  
D 4113

...

B 2359  
F 475

Name: count, Length: 6, dtype: int64

VehGas

'Regular' 10773  
'Diesel' 7977

Name: count, dtype: int64

VehBrand

B1 5069  
B2 4838

...

B11 284  
B14 136

Name: count, Length: 11, dtype: int64

Region

R24 6498  
R82 2119

...

R42 55  
R43 26

Name: count, Length: 22, dtype: int64



# Preprocess ordinal & continuous

```

1 from sklearn.compose import make_column_transformer
2
3 ct = make_column_transformer(
4     (OrdinalEncoder(), ["Area", "VehGas"]),
5     ("drop", ["VehBrand", "Region"]),
6     remainder=StandardScaler(),
7     verbose_feature_names_out=False
8 )
9 X_train_ct = ct.fit_transform(X_train)

```

```
1 X_train.head(3)
```

|   | Exposure | Area | VehPower |
|---|----------|------|----------|
| 0 | 1.00     | C    | 6        |
| 1 | 0.36     | C    | 4        |
| 2 | 0.02     | E    | 12       |

```
1 X_train_ct.head(3)
```

|   | Area | VehGas | Exposure  |
|---|------|--------|-----------|
| 0 | 2.0  | 0.0    | 1.126979  |
| 1 | 2.0  | 1.0    | -0.590896 |
| 2 | 4.0  | 1.0    | -1.503517 |



# Lecture Outline

- Entity Embedding
- **Categorical Variables & Entity Embeddings**
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale By Exposure
- Dropout



# Region column



## French Administrative Regions

Source: Nell et al. (2020), [Case Study: French Motor Third-Party Liability Claims](#), SSRN.



# One-hot encoding

```

1 oe = OneHotEncoder(sparse_output=False)
2 X_train_oh = oe.fit_transform(X_train[["Region"]])
3 X_test_oh = oe.transform(X_test[["Region"]])
4 print(list(X_train["Region"][:5]))
5 X_train_oh.head()

```

['R24', 'R93', 'R11', 'R42', 'R24']

|     | Region_R11 | Region_R21 | Region_R22 | Region_R23 | Region |
|-----|------------|------------|------------|------------|--------|
| 0   | 0.0        | 0.0        | 0.0        | 0.0        | 1.0    |
| 1   | 0.0        | 0.0        | 0.0        | 0.0        | 0.0    |
| ... | ...        | ...        | ...        | ...        | ...    |
| 3   | 0.0        | 0.0        | 0.0        | 0.0        | 0.0    |
| 4   | 0.0        | 0.0        | 0.0        | 0.0        | 1.0    |

5 rows × 22 columns

# Train on one-hot inputs

```
1 num_regions = len(oe.categories_[0])
2
3 random.seed(12)
4 model = Sequential([
5     Dense(2, input_dim=num_regions),
6     Dense(1, activation="exponential")
7 ])
8
9 model.compile(optimizer="adam", loss="poisson")
10
11 es = EarlyStopping(verbose=True)
12 hist = model.fit(X_train_oh, y_train, epochs=100, verbose=0,
13                 validation_split=0.2, callbacks=[es])
14 hist.history["val_loss"][-1]
```

Epoch 12: early stopping

0.7526934146881104





# Consider the first layer

```
1 every_category = pd.DataFrame(np.eye(num_regions), columns=oe.categories_[0])
2 every_category.head(3)
```

|   | R11 | R21 | R22 | R23 | R24 | R25 | R26 | R31 | R41 | R42 | ... | R53 | R54 | R72 | R73 | R74 | R8  |
|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 1 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 2 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | ... | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

3 rows × 22 columns

```
1 # Put this through the first layer of the model
2 X = every_category.to_numpy()
3 model.layers[0](X)
```

```
<tf.Tensor: shape=(22, 2), dtype=float32, numpy=
array([[ -0.21, -0.14],
       [ 0.21, -0.17],
       [-0.22,  0.1 ],
       [-0.83,  0.1 ],
       [-0.01, -0.66],
       [-0.65, -0.13],
       [-0.36, -0.41],
       [ 0.21, -0.03],
       [-0.93, -0.57],
       [ 0.2 , -0.41],
       [-0.43, -0.21],
       [-1.13, -0.33],
       [ 0.17, -0.68],
       [-0.88, -0.55],
       [-0.13,  0.05],
       [ 0.11,  0.  ],
       [-0.46, -0.38],
       [-0.62, -0.37],
       [-0.19, -0.28],
       [-0.22,  0.15],
       [ 0.2 ,  0.16],
       [ 0.2 ,  0.16]])
```



# The first layer

```
1 layer = model.layers[0]
2 W, b = layer.get_weights()
3 X.shape, W.shape, b.shape
```

 $((22, 22), (22, 2), (2, ))$ 
$$1 \quad X @ W + b$$

```
array([[ -0.21,  -0.14],
       [  0.21,  -0.17],
       [-0.22,   0.1 ],
       [-0.83,   0.1 ],
       [-0.01,  -0.66],
       [-0.65,  -0.13],
       [-0.36,  -0.41],
       [  0.21,  -0.03],
       [-0.93,  -0.57],
       [  0.2 ,  -0.41],
       [-0.43,  -0.21],
       [-1.13,  -0.33],
       [  0.17,  -0.68],
       [-0.88,  -0.55],
       [-0.13,   0.05],
       [  0.11,   0.   ],
       [-0.46,  -0.38],
       [-0.62,  -0.37],
       [-0.19,  -0.28],
```

 $1 \quad W + b$ 

```
array([[ -0.21,  -0.14],
       [  0.21,  -0.17],
       [-0.22,   0.1 ],
       [-0.83,   0.1 ],
       [-0.01,  -0.66],
       [-0.65,  -0.13],
       [-0.36,  -0.41],
       [  0.21,  -0.03],
       [-0.93,  -0.57],
       [  0.2 ,  -0.41],
       [-0.43,  -0.21],
       [-1.13,  -0.33],
       [  0.17,  -0.68],
       [-0.88,  -0.55],
       [-0.13,   0.05],
       [  0.11,   0.  ],
       [-0.46,  -0.38],
       [-0.62,  -0.37],
       [-0.19,  -0.28],
```



# Just a look-up operation

```
1 display(list(oe.categories_[0]))
```

```
['R11',  
'R21',  
'R22',  
'R23',  
'R24',  
'R25',  
'R26',  
'R31',  
'R41',  
'R42',  
'R43',  
'R52',  
'R53',  
'R54',  
'R72',  
'R73',  
'R74',  
'R82',  
'R83',  
...]
```

```
1 W + b
```

```
array([[ -0.21,  -0.14],  
       [  0.21,  -0.17],  
       [-0.22,   0.1 ],  
       [-0.83,   0.1 ],  
       [-0.01,  -0.66],  
       [-0.65,  -0.13],  
       [-0.36,  -0.41],  
       [  0.21,  -0.03],  
       [-0.93,  -0.57],  
       [  0.2 ,  -0.41],  
       [-0.43,  -0.21],  
       [-1.13,  -0.33],  
       [  0.17,  -0.68],  
       [-0.88,  -0.55],  
       [-0.13,   0.05],  
       [  0.11,   0.  ],  
       [-0.46,  -0.38],  
       [-0.62,  -0.37],  
       [-0.19,  -0.28],  
       ...])
```



# Turn the region into an index

```
1 oe = OrdinalEncoder()  
2 X_train_reg = oe.fit_transform(X_train[["Region"]])  
3 X_test_reg = oe.transform(X_test[["Region"]])  
4  
5 for i, reg in enumerate(oe.categories_[0][:3]):  
6     print(f"The Region value {reg} gets turned into {i}.")
```

The Region value R11 gets turned into 0.

The Region value R21 gets turned into 1.

The Region value R22 gets turned into 2.



# Embedding

```
1 from keras.layers import Embedding
2 num_regions = len(np.unique(X_train[["Region"]]))
3
4 random.seed(12)
5 model = Sequential([
6     Embedding(input_dim=num_regions, output_dim=2),
7     Dense(1, activation="exponential")
8 ])
9
10 model.compile(optimizer="adam", loss="poisson")
```



# Fitting that model

```
1 es = EarlyStopping(verbose=True)
2 hist = model.fit(X_train_reg, y_train, epochs=100, verbose=0,
3                 validation_split=0.2, callbacks=[es])
4 hist.history["val_loss"][-1]
```

Epoch 5: early stopping

0.7526668906211853

```
1 model.layers
```

```
[<Embedding name=embedding, built=True>, <Dense name=dense_6, built=True>]
```



# Keras' Embedding Layer

```
1 model.layers[0].get_weights()[0]
```

```
array([[ -0.12, -0.11],
       [ 0.03, -0.   ],
       [-0.02,  0.01],
       [-0.25, -0.14],
       [-0.28, -0.32],
       [-0.3 , -0.22],
       [-0.31, -0.28],
       [ 0.1 ,  0.07],
       [-0.61, -0.51],
       [-0.06, -0.12],
       [-0.17, -0.14],
       [-0.6 , -0.46],
       [-0.22, -0.27],
       [-0.59, -0.5 ],
       [-0.   ,  0.02],
       [ 0.07,  0.06],
       [-0.31, -0.28],
       [-0.4 , -0.34],
       [-0.16, -0.15],
       ...])
```

```
1 X_train["Region"].head(4)
```

```
0    R24
1    R93
2    R11
3    R42
Name: Region, dtype: object
```

```
1 X_sample = X_train_reg[:4].to_numpy()
2 X_sample
```

```
array([[ 4.],
       [20.],
       [ 0.],
       [ 9.]])
```

```
1 enc_tensor = model.layers[0](X_sample)
2 keras.ops.convert_to_numpy(enc_tensor).
```

```
array([[ -0.28, -0.32],
       [ 0.08,  0.03],
       [-0.12, -0.11],
       [-0.06, -0.12]], dtype=float32)
```

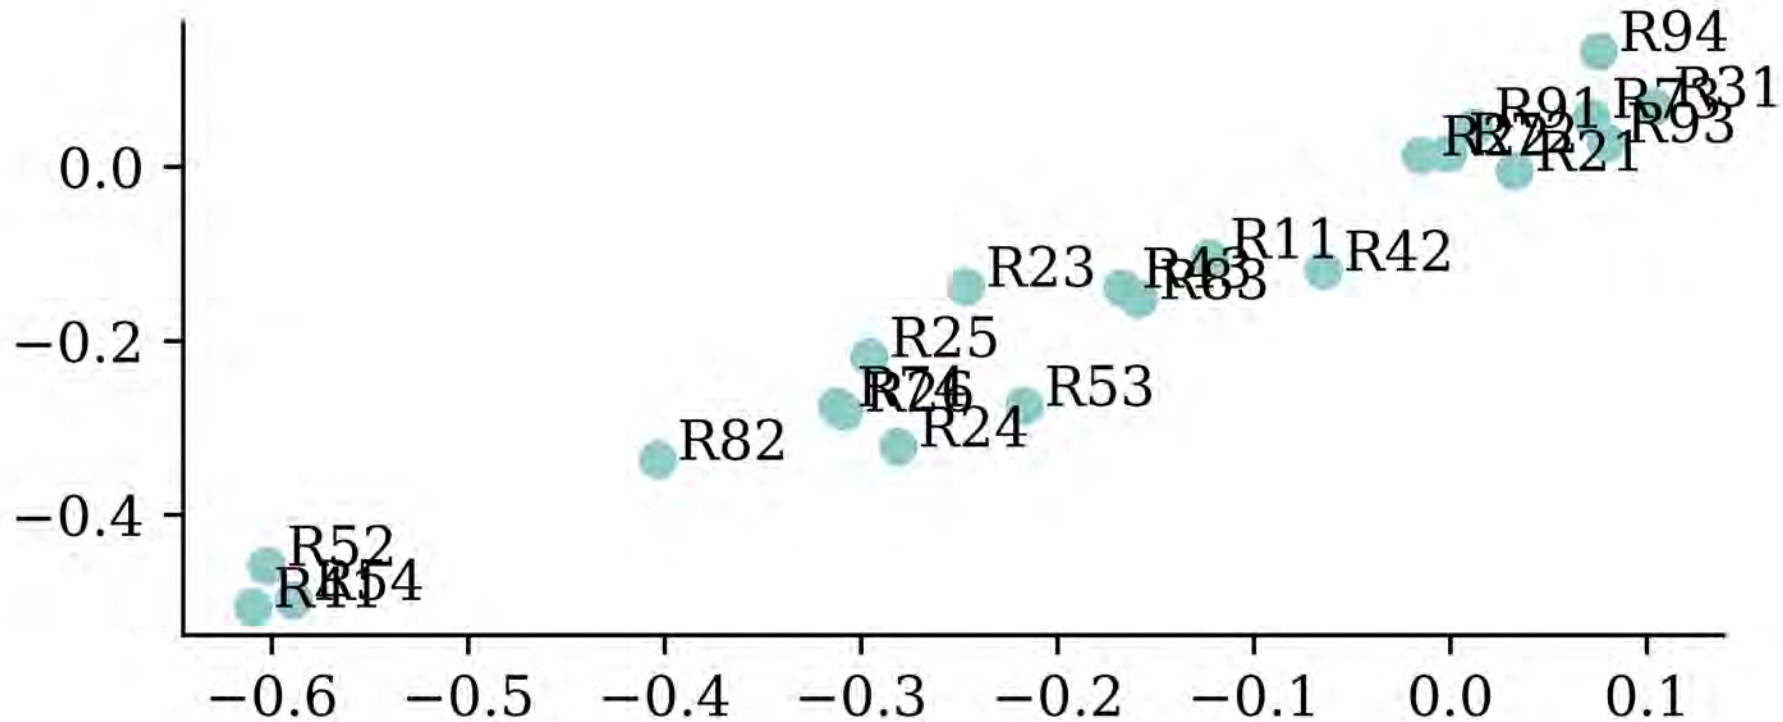


# The learned embeddings

```

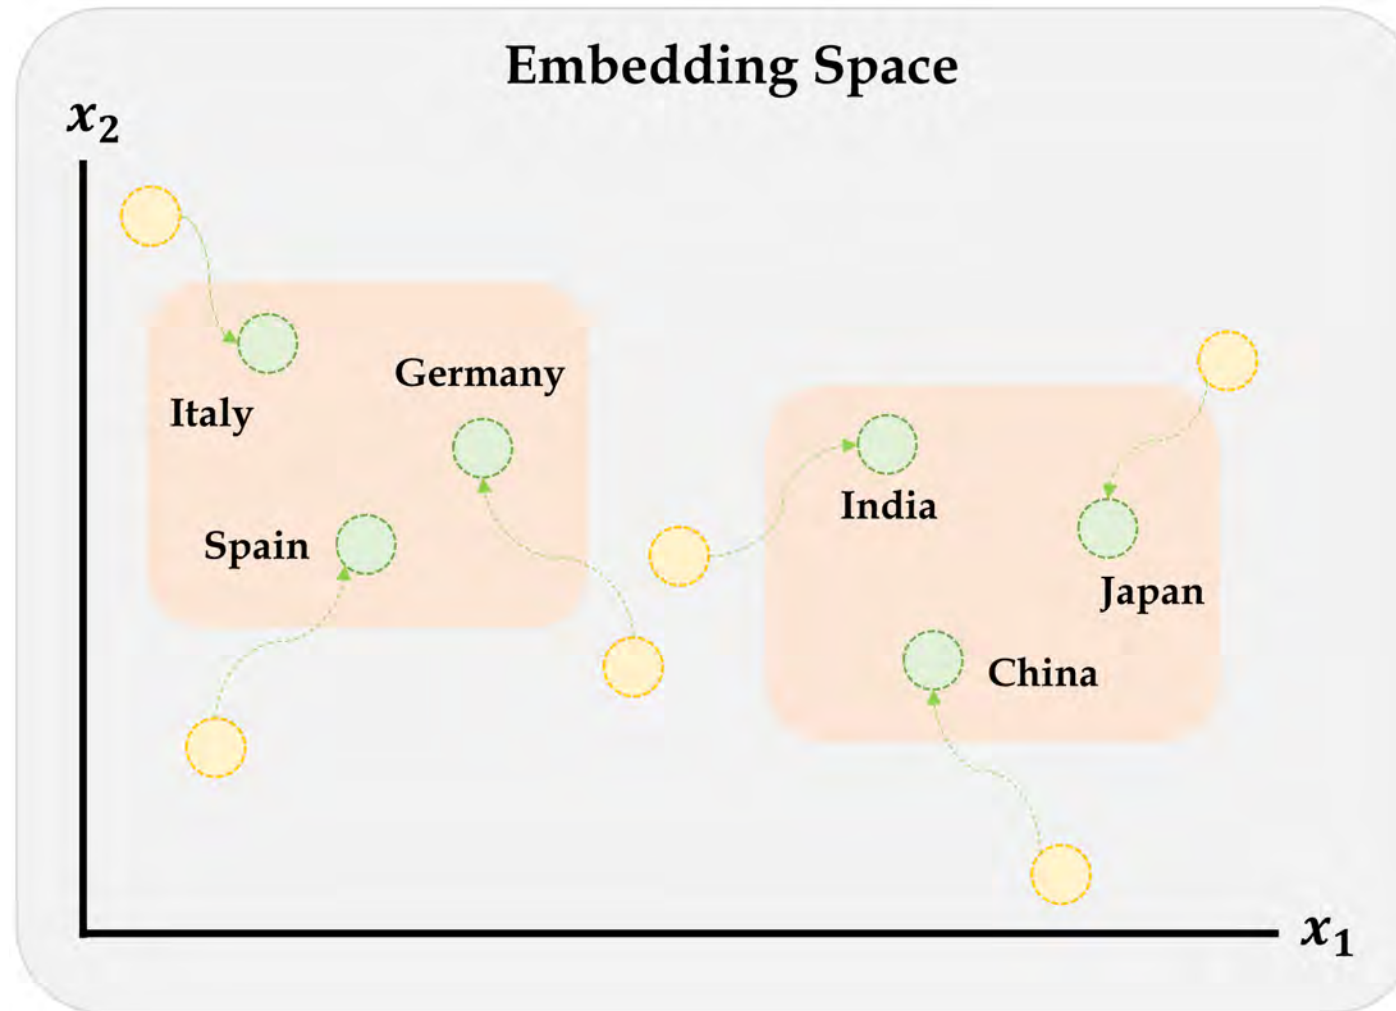
1 points = model.layers[0].get_weights()[0]
2 plt.scatter(points[:,0], points[:,1])
3 for i in range(num_regions):
4     plt.text(points[i,0]+0.01, points[i,1] , s=oe.categories_[0][i])

```





# Entity embeddings



Embeddings will gradually improve during training.

Source: Marcus Lautier (2022).



# Embeddings & other inputs

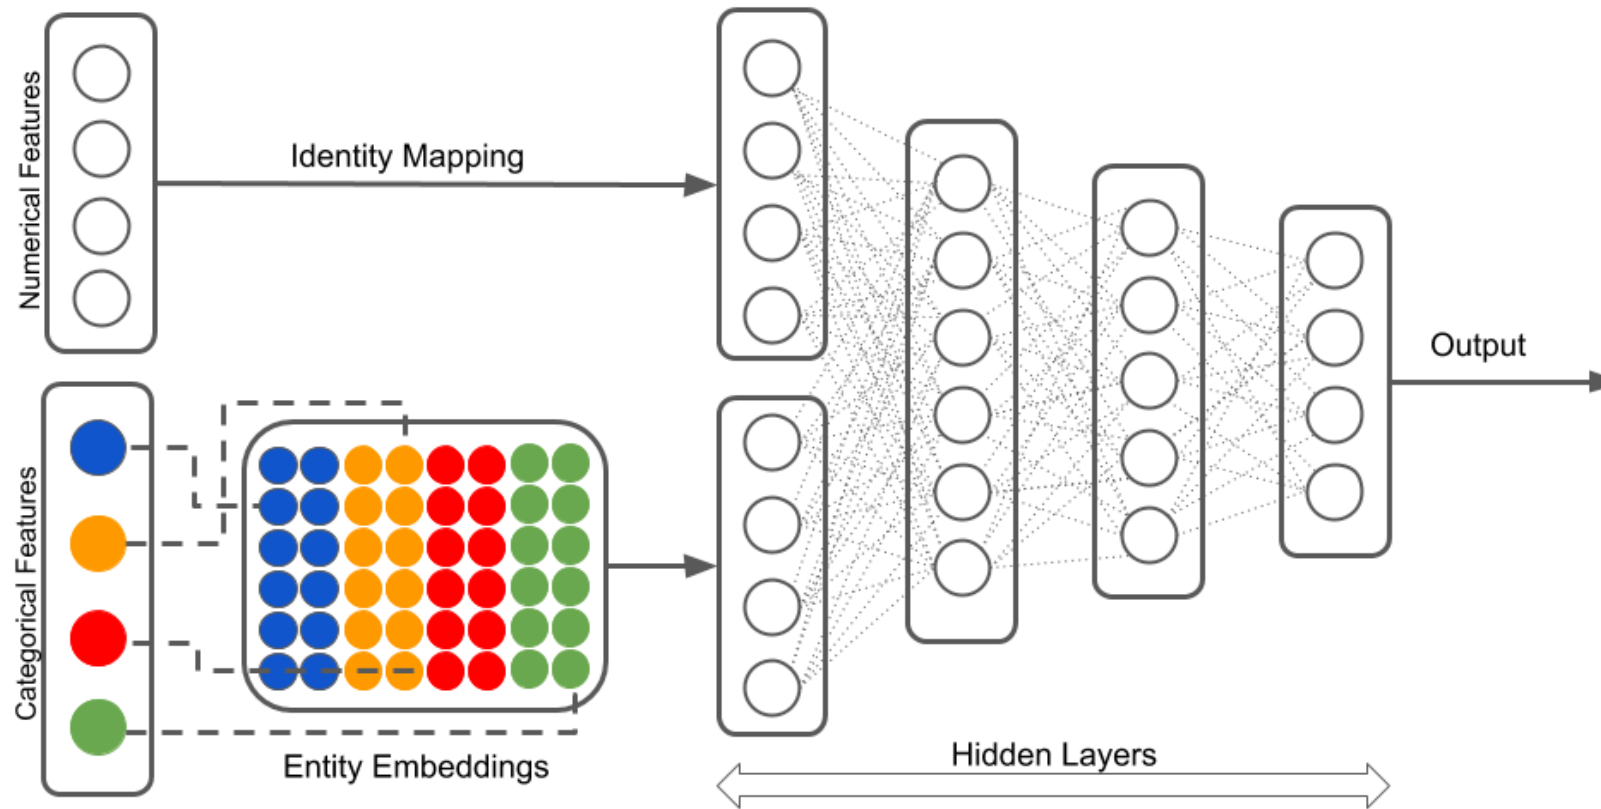


Illustration of a neural network with both continuous and categorical inputs.

We can't do this with Sequential models...

Source: LotusLabs Blog, [Accurate insurance claims prediction with Deep Learning](#).



# Lecture Outline

- Entity Embedding
- Categorical Variables & Entity Embeddings
- **Keras' Functional API**
- French Motor Dataset with Embeddings
- Scale By Exposure
- Dropout



# Converting Sequential models

```
1 from keras.models import Model
2 from keras.layers import Input
```

```
1 random.seed(12)
2
3 model = Sequential([
4     Dense(30, "leaky_relu"),
5     Dense(1, "exponential")
6 ])
7
8 model.compile(
9     optimizer="adam",
10    loss="poisson")
11
12 hist = model.fit(
13     X_train_oh, y_train,
14     epochs=1, verbose=0,
15     validation_split=0.2)
16 hist.history["val_loss"][-1]
```

0.7535399198532104

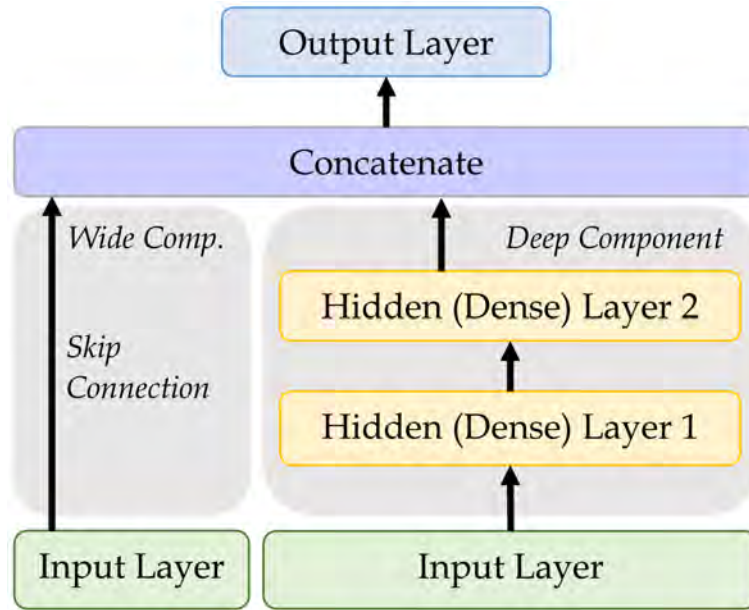
```
1 random.seed(12)
2
3 inputs = Input(shape=(X_train_oh.shape[
4 x = Dense(30, "leaky_relu")(inputs)
5 out = Dense(1, "exponential")(x)
6 model = Model(inputs, out)
7
8 model.compile(
9     optimizer="adam",
10    loss="poisson")
11
12 hist = model.fit(
13     X_train_oh, y_train,
14     epochs=1, verbose=0,
15     validation_split=0.2)
16 hist.history["val_loss"][-1]
```

0.7535399198532104

See **one-length tuples**.



# Wide & Deep network



An illustration of the wide & deep network architecture.

Add a *skip connection* from input to output layers.

```

1 from keras.layers \
2     import Concatenate
3
4 inp = Input(shape=X_train.shape[1:])
5 hidden1 = Dense(30, "leaky_relu")(inp)
6 hidden2 = Dense(30, "leaky_relu")(hidden1)
7 concat = Concatenate()(
8     [inp, hidden2])
9 output = Dense(1)(concat)
10 model = Model(
11     inputs=[inp],
12     outputs=[output])

```

# Naming the layers

For complex networks, it is often useful to give meaningful names to the layers.

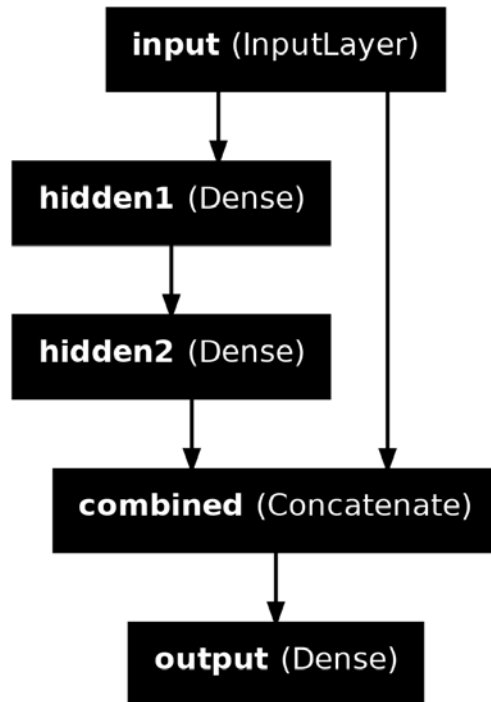
```
1 input_ = Input(shape=X_train.shape[1:], name="input")
2 hidden1 = Dense(30, activation="leaky_relu", name="hidden1")(input_)
3 hidden2 = Dense(30, activation="leaky_relu", name="hidden2")(hidden1)
4 concat = Concatenate(name="combined")([input_, hidden2])
5 output = Dense(1, name="output")(concat)
6 model = Model(inputs=[input_], outputs=[output])
```



# Inspecting a complex model

```
1 from keras.utils import plot_model
```

```
1 plot_model(model, show_shapes=True)
```



```
1 model.summary(line_length=75)
```

Model: "functional\_8"

| Layer (type)           | Output Shape | Param # | Connected to                 |
|------------------------|--------------|---------|------------------------------|
| input (InputLayer)     | (None, 10)   | 0       | -                            |
| hidden1 (Dense)        | (None, 30)   | 330     | input[0][0]                  |
| hidden2 (Dense)        | (None, 30)   | 930     | hidden1[0][0]                |
| combined (Concatenate) | (None, 40)   | 0       | input[0][0]<br>hidden2[0][0] |
| output (Dense)         | (None, 1)    | 41      | combined[0][0]               |

Total params: 1,301 (5.08 KB)  
 Trainable params: 1,301 (5.08 KB)  
 Non-trainable params: 0 (0.00 B)

# Lecture Outline

- Entity Embedding
- Categorical Variables & Entity Embeddings
- Keras' Functional API
- **French Motor Dataset with Embeddings**
- Scale By Exposure
- Dropout





# The desired architecture

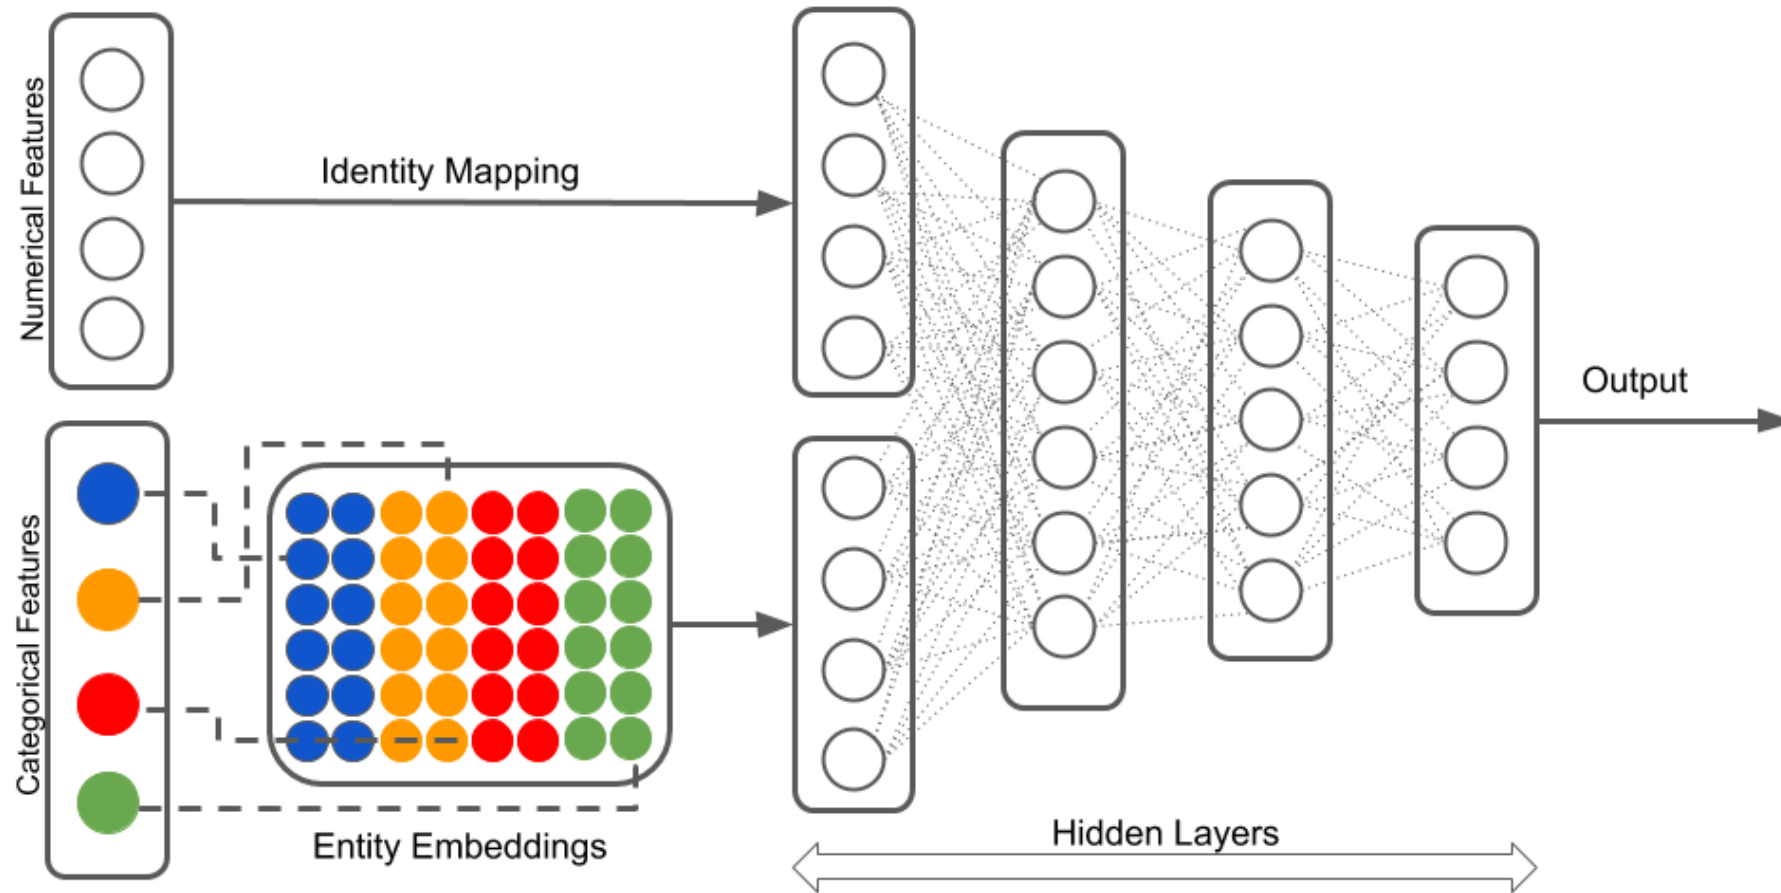


Illustration of a neural network with both continuous and categorical inputs.

Source: LotusLabs Blog, [Accurate insurance claims prediction with Deep Learning](#).



# Preprocess all French motor inputs

Transform the categorical variables to integers:

```
1 num_brands, num_regions = X_train.nunique()[["VehBrand", "Region"]]
2
3 ct = make_column_transformer(
4     (OrdinalEncoder(), ["VehBrand", "Region", "Area", "VehGas"]),
5     remainder=StandardScaler(),
6     verbose_feature_names_out=False
7 )
8 X_train_ct = ct.fit_transform(X_train)
9 X_test_ct = ct.transform(X_test)
```

Split the brand and region data apart from the rest:

```
1 X_train_brand = X_train_ct["VehBrand"]; X_test_brand = X_test_ct["VehBrand"]
2 X_train_region = X_train_ct["Region"]; X_test_region = X_test_ct["Region"]
3 X_train_rest = X_train_ct.drop(["VehBrand", "Region"], axis=1)
4 X_test_rest = X_test_ct.drop(["VehBrand", "Region"], axis=1)
```



# Organise the inputs

Make a Keras **Input** for: vehicle brand, region, & others.

```
1 veh_brand = Input(shape=(1,), name="vehBrand")
2 region = Input(shape=(1,), name="region")
3 other_inputs = Input(shape=X_train_rest.shape[1:], name="otherInputs")
```

Create embeddings and join them with the other inputs.

```
1 from keras.layers import Reshape
2
3 random.seed(1337)
4 veh_brand_ee = Embedding(input_dim=num_brands, output_dim=2,
5     name="vehBrandEE")(veh_brand)
6 veh_brand_ee = Reshape(target_shape=(2,))(veh_brand_ee)
7
8 region_ee = Embedding(input_dim=num_regions, output_dim=2,
9     name="regionEE")(region)
10 region_ee = Reshape(target_shape=(2,))(region_ee)
11
12 x = Concatenate(name="combined")([veh_brand_ee, region_ee, other_inputs])
```



# Complete the model and fit it

Feed the combined embeddings & continuous inputs to some normal dense layers.

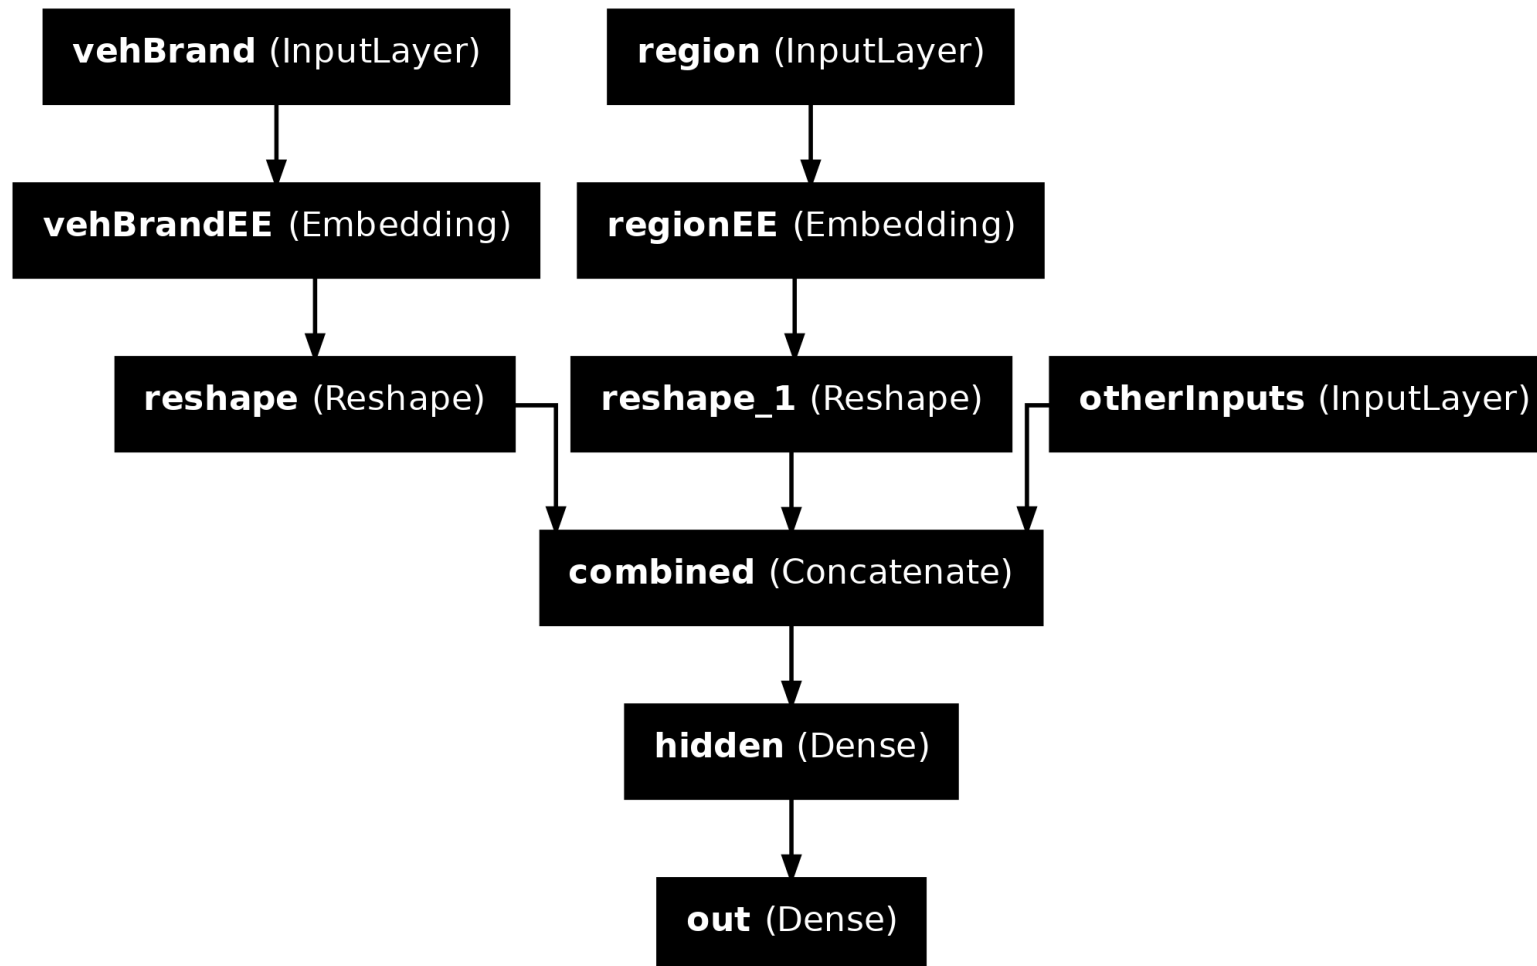
```
1 x = Dense(30, "relu", name="hidden")(x)
2 out = Dense(1, "exponential", name="out")(x)
3
4 model = Model([veh_brand, region, other_inputs], out)
5 model.compile(optimizer="adam", loss="poisson")
6
7 hist = model.fit((X_train_brand, X_train_region, X_train_rest),
8                 y_train, epochs=100, verbose=0,
9                 callbacks=[EarlyStopping(patience=5)], validation_split=0.2)
10 np.min(hist.history["val_loss"])
```

0.6692155599594116



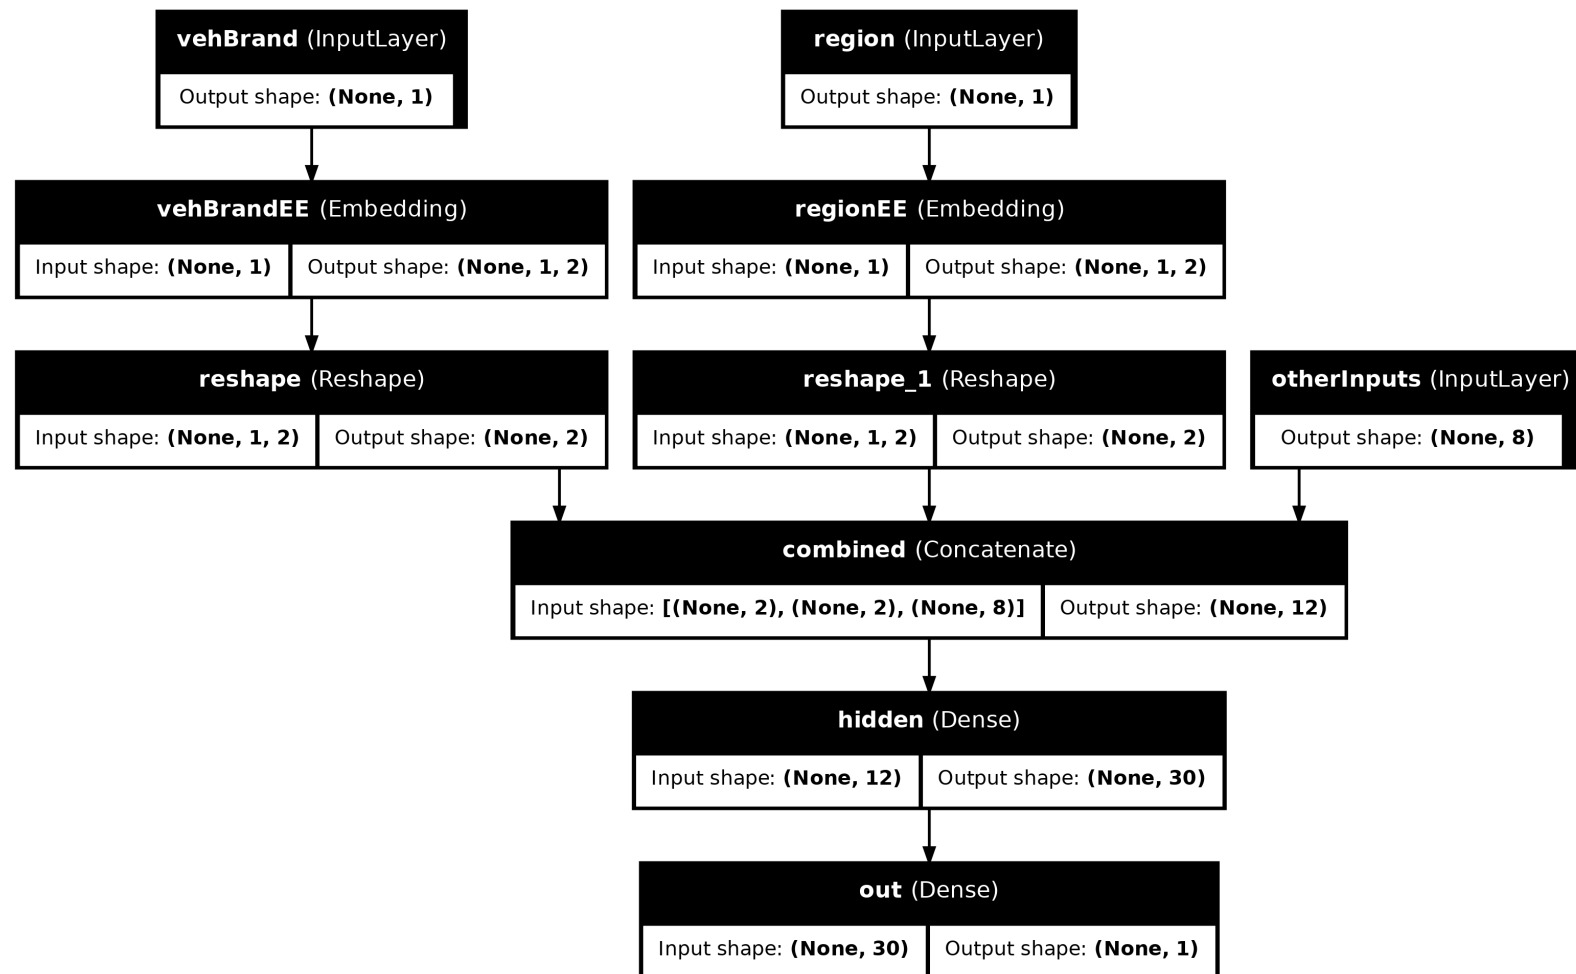
# Plotting this model

```
1 plot_model(model, show_layer_names=True)
```



# Why we need to reshape

```
1 plot_model(model, show_layer_names=True, show_shapes=True)
```



# Lecture Outline

- Entity Embedding
- Categorical Variables & Entity Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- **Scale By Exposure**
- Dropout



# Two different models

Have  $\{(\mathbf{x}_i, y_i)\}_{i=1, \dots, n}$  for  $\mathbf{x}_i \in \mathbb{R}^{47}$  and  $y_i \in \mathbb{N}_0$ .

**Model 1:** Say  $Y_i \sim \text{Poisson}(\lambda(\mathbf{x}_i))$ .

But, the exposures are different for each policy.  $\lambda(\mathbf{x}_i)$  is the expected number of claims for the duration of policy  $i$ 's contract.

**Model 2:** Say  $Y_i \sim \text{Poisson}(\text{Exposure}_i \times \lambda(\mathbf{x}_i))$ .

Now,  $\text{Exposure}_i \notin \mathbf{x}_i$ , and  $\lambda(\mathbf{x}_i)$  is the rate *per year*.





# Just take continuous variables

```

1 ct = make_column_transformer(
2     ("passthrough", ["Exposure"]),
3     ("drop", ["VehBrand", "Region", "Area", "VehGas"]),
4     remainder=StandardScaler(),
5     verbose_feature_names_out=False
6 )
7 X_train_ct = ct.fit_transform(X_train)
8 X_test_ct = ct.transform(X_test)

```

Split exposure apart from the rest:

```

1 X_train_exp = X_train_ct["Exposure"]; X_test_exp = X_test_ct["Exposure"]
2 X_train_rest = X_train_ct.drop("Exposure", axis=1)
3 X_test_rest = X_test_ct.drop("Exposure", axis=1)

```

Organise the inputs:

```

1 exposure = Input(shape=(1,), name="exposure")
2 other_inputs = Input(shape=X_train_rest.shape[1:], name="otherInputs")

```



# Make & fit the model

Feed the continuous inputs to some normal dense layers.

```
1 random.seed(1337)
2 x = Dense(30, "relu", name="hidden1")(other_inputs)
3 x = Dense(30, "relu", name="hidden2")(x)
4 lambda_ = Dense(1, "exponential", name="lambda")(x)
```

```
1 from keras.layers import Multiply
2
3 out = Multiply(name="out")([lambda_, exposure])
4 model = Model([exposure, other_inputs], out)
5 model.compile(optimizer="adam", loss="poisson")
6
7 es = EarlyStopping(patience=10, restore_best_weights=True, verbose=1)
8 hist = model.fit((X_train_exp, X_train_rest),
9                 y_train, epochs=100, verbose=0,
10                callbacks=[es], validation_split=0.2)
11 np.min(hist.history["val_loss"])
```

Epoch 40: early stopping

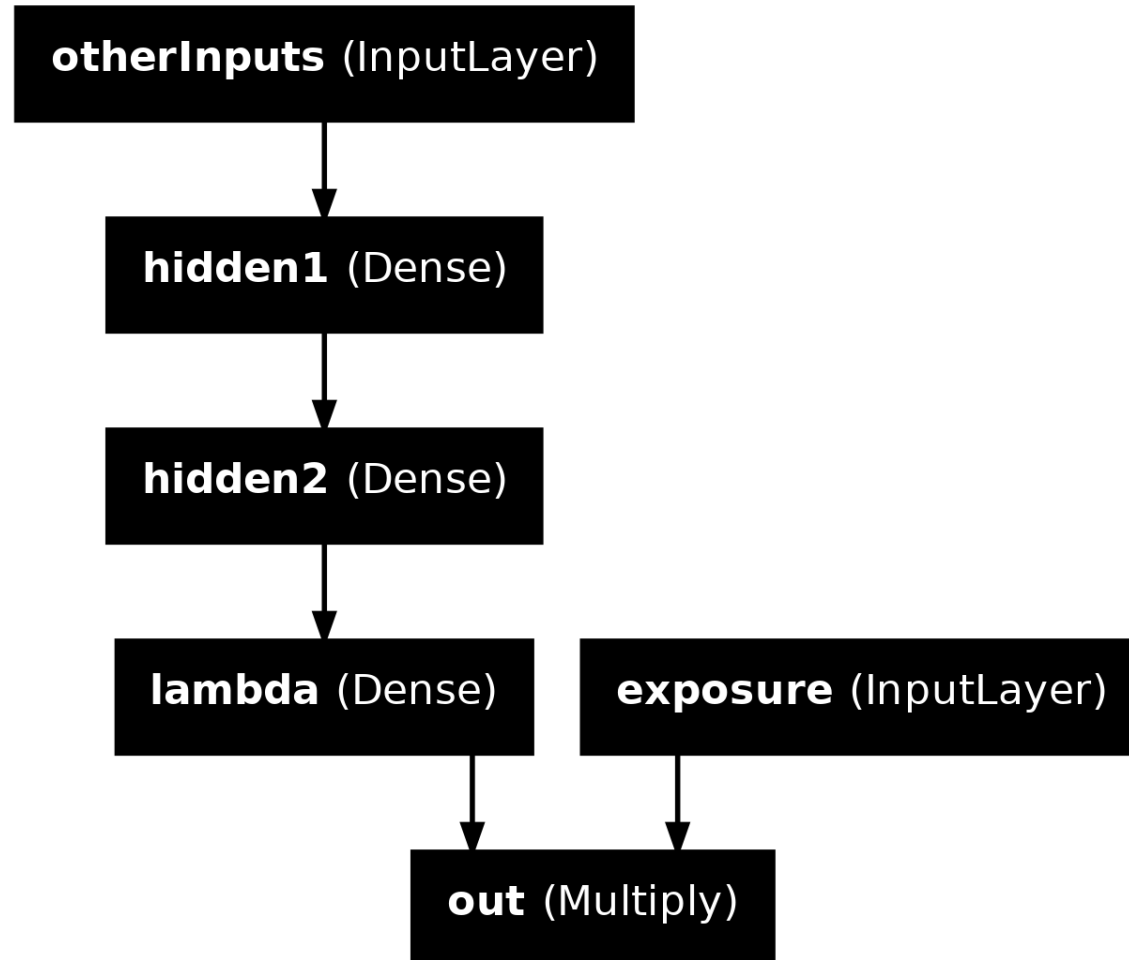
Restoring model weights from the end of the best epoch: 30.

0.8829042911529541



# Plot the model

```
1 plot_model(model, show_layer_names=True)
```

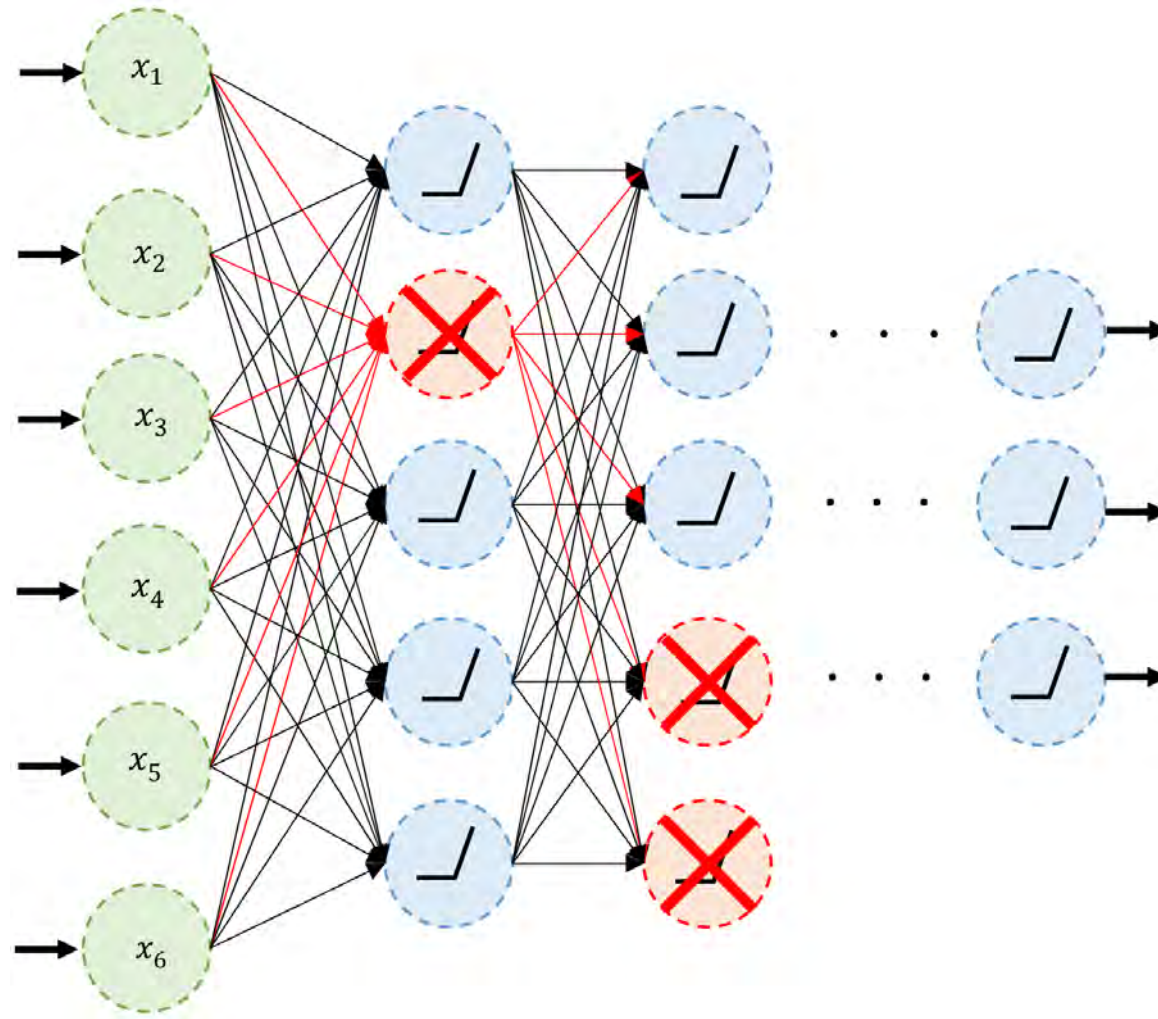


# Lecture Outline

- Entity Embedding
- Categorical Variables & Entity Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale By Exposure
- **Dropout**



# Dropout



An example of neurons dropped during training.

Sources: Marcus Lautier (2022).



# Dropout quote #1

It's surprising at first that this destructive technique works at all. Would a company perform better if its employees were told to toss a coin every morning to decide whether or not to go to work? Well, who knows; perhaps it would! The company would be forced to adapt its organization; it could not rely on any single person to work the coffee machine or perform any other critical tasks, so this expertise would have to be spread across several people. Employees would have to learn to cooperate with many of their coworkers, not just a handful of them.



## Dropout quote #2

The company would become much more resilient. If one person quit, it wouldn't make much of a difference. It's unclear whether this idea would actually work for companies, but it certainly does for neural networks. Neurons trained with dropout cannot co-adapt with their neighboring neurons; they have to be as useful as possible on their own. They also cannot rely excessively on just a few input neurons; they must pay attention to each of their input neurons. They end up being less sensitive to slight changes in the inputs. In the end, you get a more robust network that generalizes better.



# Code: Dropout

Dropout is just another layer in Keras.

```
1 from keras.layers import Dropout
2
3 random.seed(2);
4
5 model = Sequential([
6     Dense(30, activation="leaky_relu", name="hidden1"),
7     Dropout(0.2),
8     Dense(30, activation="leaky_relu", name="hidden2"),
9     Dropout(0.2),
10    Dense(1, activation="exponential", name="output")
11 ])
12
13 model.compile("adam", "mse")
14 model.fit(X_train_ct, y_train, epochs=4, verbose=0);
```





# Code: Dropout after training

Making predictions is the same as any other model:

```
1 model.predict(X_train_ct.head(3),
2               verbose=0)
```

```
array([[0.64],
       [0.14],
       [0.46]], dtype=float32)
```

```
1 model.predict(X_train_ct.head(3),
2               verbose=0)
```

```
array([[0.64],
       [0.14],
       [0.46]], dtype=float32)
```

We can make the model think it is still training:

```
1 model(X_train_ct.head(3).to_numpy(),
2       training=True).numpy()
```

```
array([[0.64],
       [0.11],
       [0.48]], dtype=float32)
```

```
1 model(X_train_ct.head(3).to_numpy(),
2       training=True).numpy()
```

```
array([[0.68],
       [0.16],
       [0.55]], dtype=float32)
```



# Dropout Limitation

- Increased Training Time: Since dropout introduces noise into the training process, it can make the training process slower.
- Sensitivity to Dropout Rates: the performance of dropout is highly dependent on the chosen dropout rate.



# Deep Ensembles

It's simple to implement and requires very little hyperparameter tuning.

We summarise the deep ensemble approach for uncertainty quantification as follows:

1. Train  $D$  neural networks with different random weights initialisations independently in parallel. The trained weights are  $\mathbf{w}^{(1)}, \dots, \mathbf{w}^{(D)}$ .

## Note

Unfinished section... Please imagine some wildly interesting content here.



# Package Versions

```
1 from watermark import watermark
2 print(watermark(python=True, packages="keras,matplotlib,numpy,pandas,seaborn,scipy,torch"))
```

```
Python implementation: CPython
Python version       : 3.11.9
IPython version      : 8.24.0
```

```
keras      : 3.3.3
matplotlib: 3.9.0
numpy      : 1.26.4
pandas     : 2.2.2
seaborn    : 0.13.2
scipy      : 1.11.0
torch      : 2.3.1
tensorflow: 2.16.1
tf_keras   : 2.16.0
```



# Glossary

- dropout
- entity embeddings
- Input layer
- Keras functional API
- Reshape layer
- skip connection
- wide & deep network structure

